

# hw1\_yc

February 6, 2024

```
[ ]: import torch
      from torch import nn
      import pandas as pd
      import numpy as np
      from torch.utils.data import DataLoader, TensorDataset
      from torch.nn.functional import mse_loss
```

```
[ ]: class PoissonRegression(nn.Module):
      def __init__(self, in_dim):
          super().__init__()
          self.w = nn.Linear(in_dim, 1)
          nn.init.normal_(self.w.weight, 0, .01)
          nn.init.constant_(self.w.bias, 0)

      def forward(self, x):
          return torch.exp(self.w(x))

      def poisson_loss(y_hat, y):
          loss = torch.mean(y_hat - y * torch.log(y_hat))
          return loss
```

```
[ ]: df = pd.read_csv("HW1_poisson_regression_data.csv")

      X = df[['X1', 'X2', 'X3', 'X4', 'X5']].values
      y = df['Y'].values

      X_tensor = torch.tensor(X, dtype=torch.float)
      y_tensor = torch.tensor(y, dtype=torch.float)

      df.head(n=10)
```

```
[ ]:      Y      X1      X2      X3      X4      X5
0    5  0.422447  0.971883  0.533636  0.058229  0.912143
1    2  0.382487  0.507943  0.950425  0.174120  0.551095
2    4  0.742815  0.789894  0.374025  0.320585  0.286143
3    6  0.905309  0.093778  0.552340  0.909098  0.836523
4    7  0.557285  0.567502  0.552205  0.988464  0.989089
```

5	3	0.708192	0.949948	0.151380	0.110029	0.491539
6	2	0.547922	0.835359	0.242528	0.908562	0.180607
7	1	0.262783	0.214829	0.468023	0.839837	0.004941
8	13	0.886888	0.567713	0.657572	0.311464	0.879461
9	3	0.279256	0.168646	0.441679	0.164778	0.623770

```
[ ]: model = PoissonRegression(in_dim=5)

optimizer = torch.optim.SGD(model.parameters(), lr=0.001)

max_epochs = 1000
batch_size = 32

dataset = TensorDataset(X_tensor, y_tensor)
dataloader = DataLoader(dataset, batch_size=batch_size, shuffle=True)

# lists to store loss and MSE values
loss_history = []
mse_history = []

for epoch in range(max_epochs):

    for x_batch, y_batch in dataloader:
        y_hat = model(x_batch)
        loss = poisson_loss(y_hat, y_batch)
        mse = mse_loss(y_hat, y_batch.unsqueeze(1), reduction='mean')

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    loss_history.append(loss.item())
    mse_history.append(mse.item())

    if epoch % 20 == 0:
        print(f"Epoch {epoch}, Loss: {loss.item()}, MSE: {mse:.4f}")

# After training, check model parameters
for name, param in model.named_parameters():
    if param.requires_grad:
        print(f"{name}: {param.data}")
```

```
Epoch 0, Loss: 0.3592326045036316, MSE: 16.9577
Epoch 20, Loss: -4.5003156661987305, MSE: 36.6645
Epoch 40, Loss: -3.210230588912964, MSE: 19.1162
Epoch 60, Loss: -2.2181508541107178, MSE: 10.2717
Epoch 80, Loss: -4.463136196136475, MSE: 27.1496
```

Epoch 100, Loss: -1.152472734451294, MSE: 6.8575  
 Epoch 120, Loss: -5.613916397094727, MSE: 66.0387  
 Epoch 140, Loss: -5.388947486877441, MSE: 79.0067  
 Epoch 160, Loss: -0.12387394905090332, MSE: 15.8548  
 Epoch 180, Loss: 0.16124770045280457, MSE: 5.4298  
 Epoch 200, Loss: 1.0899252891540527, MSE: 9.7699  
 Epoch 220, Loss: -0.6478701829910278, MSE: 9.9053  
 Epoch 240, Loss: -2.731973886489868, MSE: 13.7152  
 Epoch 260, Loss: -3.1296634674072266, MSE: 9.5168  
 Epoch 280, Loss: -3.638622283935547, MSE: 17.1657  
 Epoch 300, Loss: -0.8484793305397034, MSE: 5.1116  
 Epoch 320, Loss: -2.0533597469329834, MSE: 7.4563  
 Epoch 340, Loss: 1.74629545211792, MSE: 8.5326  
 Epoch 360, Loss: 0.9207884073257446, MSE: 6.5165  
 Epoch 380, Loss: -4.037039756774902, MSE: 25.7735  
 Epoch 400, Loss: -4.262198448181152, MSE: 12.3523  
 Epoch 420, Loss: -6.347617149353027, MSE: 66.7400  
 Epoch 440, Loss: -1.894134283065796, MSE: 17.0587  
 Epoch 460, Loss: 2.7959353923797607, MSE: 10.2595  
 Epoch 480, Loss: -0.6574552655220032, MSE: 5.2082  
 Epoch 500, Loss: -3.8678395748138428, MSE: 18.9355  
 Epoch 520, Loss: -1.3649650812149048, MSE: 21.5800  
 Epoch 540, Loss: -4.378934383392334, MSE: 80.6029  
 Epoch 560, Loss: -2.0753934383392334, MSE: 7.5090  
 Epoch 580, Loss: 0.21774227917194366, MSE: 3.9014  
 Epoch 600, Loss: -4.617292404174805, MSE: 49.5742  
 Epoch 620, Loss: 0.4108962416648865, MSE: 5.7028  
 Epoch 640, Loss: 0.060598939657211304, MSE: 7.0041  
 Epoch 660, Loss: -0.6565761566162109, MSE: 9.9542  
 Epoch 680, Loss: -5.871087551116943, MSE: 41.9801  
 Epoch 700, Loss: -0.6556139588356018, MSE: 4.9551  
 Epoch 720, Loss: -2.2558822631835938, MSE: 14.1363  
 Epoch 740, Loss: 1.658700942993164, MSE: 8.2224  
 Epoch 760, Loss: -1.3634697198867798, MSE: 7.0020  
 Epoch 780, Loss: 0.22795042395591736, MSE: 5.0803  
 Epoch 800, Loss: -2.430828094482422, MSE: 50.3383  
 Epoch 820, Loss: -6.403717041015625, MSE: 39.6533  
 Epoch 840, Loss: -2.96830153465271, MSE: 20.5044  
 Epoch 860, Loss: -3.5061047077178955, MSE: 16.2696  
 Epoch 880, Loss: -3.498863935470581, MSE: 21.7399  
 Epoch 900, Loss: -1.183661699295044, MSE: 4.7062  
 Epoch 920, Loss: -0.3022191524505615, MSE: 5.4294  
 Epoch 940, Loss: 1.2656785249710083, MSE: 7.9844  
 Epoch 960, Loss: -0.4763941764831543, MSE: 8.7847  
 Epoch 980, Loss: -2.788625955581665, MSE: 10.9425  
 w.weight: tensor([[ 0.0742, -0.0277, 0.0288, -0.0120, 0.0156]])  
 w.bias: tensor([1.3839])

```
[ ]: import matplotlib.pyplot as plt

# Plotting the losses
plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(loss_history, label='Loss')
plt.title('Loss over Epochs')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(mse_history, label='MSE')
plt.title('MSE over Epochs')
plt.xlabel('Epoch')
plt.ylabel('MSE')
plt.legend()

plt.tight_layout()
plt.show()
```

